

# Heart Disease Prediction using Support Vector Machine (SVM)

## 1. Introduction:

Data modelling is the process of identifying the patterns and discovering new insights from large dataset by using statistical methods and machine learning algorithms. There are various machine learning algorithms which are broadly classified into 2 groups - Supervised and Unsupervised learning models based on whether the output variable is labelled or not. In Supervised learning model, if the output variable is discrete or categorical, we should use the classification algorithm like decision tree or random forest to train the model. If the output variable is numeric, we should use the regression algorithm to train the model. Similarly, for Unsupervised learning model, if the output variable is not labelled, K- mean clustering algorithm should be used to divide the dataset into small clusters. The number and types of the parameter we pass while training a model also plays an important role in choosing the right algorithm. Support Vector Machine algorithms is preferred when multiple parameters are required to train the model.

The Support Vector Machine is a classification algorithm used to separate two classes of the target data by a hyperplane that has maximum margin, i.e the maximum distance between data points of each class. These data points that determines the position and the orientation of the hyperplane are called support vectors. SVM works well for high-dimensional dataset where the number of features/variables is greater than the number of observations/records.

In this project, we will analyse heart disease dataset which has 14 variables and 300 records. Since variable not named, it is hard to find the target and prediction variables. All the variables are of integer type except for the variable 10 which is a numeric. As per the instruction, the first 13 variables will be used for predicting the 14<sup>th</sup> variable which is a target variable for the given dataset. The values in the first 13 columns looks to be a information of a patient which is used to predict the binary value(0 and 1)in the 14<sup>th</sup> column, where 0 represents the absence of the heart disease and the 1 represents the presence of the heart disease.

In order to better understand the dataset, I looked for the similar dataset in Kaggle.com and found the one in the below link to be matching with our heart disease dataset.

## 2. Methods & analysis

### 2.1. Libraries used

The caret package (Classification and Regression Training) contains all necessary functions for building and training complex classification and regression problems like data splitting, pre-processing, feature selection, tuning, etc.,.

### 2.2. Loading the dataset

The heart\_tidy.csv dataset which we used for this project is a comma separated values file. We have used setwd() function to set the working directory to the path of the dataset location. We then use read.csv() function with filename as the parameter to load the csv file as a dataframe into a variable *heart\_df*. The sep parameter tell the interprets which delimiter is used as a separator in the csv file. In our dataset, since the 1<sup>st</sup> column is not the header, we have passed the value FALSE for the parameter header.

## 3. Exploratory Data Analysis

### 3.1. Structure of the dataset

We use the str() function to check the structure of the *heart\_df* dataframe.

```
> str(heart_df)
'data.frame': 300 obs. of 14 variables:
 $ v1 : int 63 67 67 37 41 56 62 57 63 53 ...
 $ v2 : int 1 1 1 1 0 1 0 0 1 1 ...
 $ v3 : int 1 4 4 3 2 2 4 4 4 4 ...
 $ v4 : int 145 160 120 130 130 120 140 120 130 140 ...
 $ v5 : int 233 286 229 250 204 236 268 354 254 203 ...
 $ v6 : int 1 0 0 0 0 0 0 0 0 1 ...
 $ v7 : int 2 2 2 0 2 0 2 0 2 2 ...
 $ v8 : int 150 108 129 187 172 178 160 163 147 155 ...
 $ v9 : int 0 1 1 0 0 0 0 1 0 1 ...
 $ v10: num 2.3 1.5 2.6 3.5 1.4 0.8 3.6 0.6 1.4 3.1 ...
 $ v11: int 3 2 2 3 1 1 3 1 2 3 ...
 $ v12: int 0 3 2 0 0 0 2 0 1 0 ...
 $ v13: int 6 3 7 3 3 3 3 3 7 7 ...
 $ v14: int 0 1 1 0 0 0 1 0 1 1 ...
```

Figure 1: Represents the structure of the *heart\_df* dataframe

From fig.1, we can understand that the dataframe has 300 observations and 14 variables. Out of 14 variables, 13 are of integer type and the variable 10 is numeric type. Since, *V14* is our target variable, we can use remaining 13 variables as an input or predictor variable. We can observe that the values of the *V14* variable is binary (0 and 1).

We use the `head()` function to check the first few rows of the `heart_df` dataframe.

```
> head(heart_df)
  V1 V2 V3 V4 V5 V6 V7 V8 V9 V10 V11 V12 V13 V14
1 63  1  1 145 233  1  2 150  0 2.3  3  0  6  0
2 67  1  4 160 286  0  2 108  1 1.5  2  3  3  1
3 67  1  4 120 229  0  2 129  1 2.6  2  2  7  1
4 37  1  3 130 250  0  0 187  0 3.5  3  0  3  0
5 41  0  2 130 204  0  2 172  0 1.4  1  0  3  0
6 56  1  2 120 236  0  0 178  0 0.8  1  0  3  0
```

Figure 2: Represents the first few rows of the `heart_df` dataframe

From fig.2, we can understand that except for V10 whose values are in decimal, the values of the rest all columns are integers.

### 3.2 Data Splicing

Data Splicing is the process of splitting the data into a training set and a testing set. We will split the data using two random samples without replacement:

- training dataset (70%) for model building
- test dataset (30%) to validate the efficiency of the model.

We have used the `set.seed()` function with the argument 3033 to get the reproducible random numbers. During data partitioning, if the same value is passed in the `set.seed()` method, we will get the identical split of training and testing data even though it is random.

For partitioning the data into training and testing set, we use the function `createDataPartition()` which is part of the `caret` package. The first parameter 'y' takes the value of the target variable which needs to be partitioned. We have passed `heart_df$V14`, the 14<sup>th</sup> column on our dataframe as it's our target variable.

The "p" parameter indicates the percentage of the split between the training and testing data and it holds a decimal value in the range of 0-1. Since we want to split the training and testing data in 70:30 ratio, we have given `p=0.7`. Out of 300 observations from the `heart_df` dataframe, 70% of records (210 observations) are spitted and loaded into a training data frame. The value `P= 0.7` determines that 70% of data should be loaded into the training data frame and the remaining 30% loaded into the testing data frame.

The list parameter returns the output as list if we pass the value TRUE and it returns the output as matrix if we pass the value as FALSE. Since we have pass FALSE in the list parameter, the createDataPartition() function returning a matrix “intrain”.

We pass the values of the “intrain” matrix which holds the indices of the observation the heart\_df dataframe to split it into training and testing data.

We use the dim() function to check the dimension of the training and testing data

```
> dim(training); dim(testing);  
[1] 210 14  
[1] 90 14
```

Figure 3: Represents the dimension of the training and testing data

From fig.3, we can understand that the training dataset has 210 observations and 14 variables whereas the testing dataset has 90 observations and 14 variables.

### 3.3 Finding the missing values in the dataset

In general, the data may contain null values, missing values, unwanted data, outliers that may lead to wrong interpretations. Hence data pre-processing is a critical step where in the irrelevant data can be removed so we can achieve better accuracy while predicting the values using our model. We use anyNA() function to check the presence of any NA(Not Available) values in the heart\_df dataframe.

```
> anyNA(heart_df)  
[1] FALSE
```

Figure 4: The output of the anyNA() function.

The output FALSE in the fig.4 shows there is no missing values in the heart\_df dataframe

### 3.3 Dataset summarized details

We use the summary() function to check the summarized details of our dataframe. It gives the descriptive statistical values such as mean, median, minimum, maximum, 1<sup>st</sup> and 3<sup>rd</sup> quantile of each variable.

```

> summary(heart_df)
      v1      v2      v3      v4      v5
Min.   :29.00  Min.   :0.00  Min.   :1.000  Min.   : 94.0  Min.   :126.0
1st Qu.:48.00  1st Qu.:0.00  1st Qu.:3.000  1st Qu.:120.0  1st Qu.:211.0
Median :56.00  Median :1.00  Median :3.000  Median :130.0  Median :241.5
Mean   :54.48  Mean   :0.68  Mean   :3.153  Mean   :131.6  Mean   :246.9
3rd Qu.:61.00  3rd Qu.:1.00  3rd Qu.:4.000  3rd Qu.:140.0  3rd Qu.:275.2
Max.   :77.00  Max.   :1.00  Max.   :4.000  Max.   :200.0  Max.   :564.0
      v6      v7      v8      v9     v10
Min.   :0.0000  Min.   :0.0000  Min.   : 71.0  Min.   :0.0000  Min.   :0.00
1st Qu.:0.0000  1st Qu.:0.0000  1st Qu.:133.8  1st Qu.:0.0000  1st Qu.:0.00
Median :0.0000  Median :0.5000  Median :153.0  Median :0.0000  Median :0.80
Mean   :0.1467  Mean   :0.9867  Mean   :149.7  Mean   :0.3267  Mean   :1.05
3rd Qu.:0.0000  3rd Qu.:2.0000  3rd Qu.:166.0  3rd Qu.:1.0000  3rd Qu.:1.60
Max.   :1.0000  Max.   :2.0000  Max.   :202.0  Max.   :1.0000  Max.   :6.20
      v11     v12     v13     v14
Min.   :1.000  Min.   :0.00  Min.   :3.000  Min.   :0.00
1st Qu.:1.000  1st Qu.:0.00  1st Qu.:3.000  1st Qu.:0.00
Median :2.000  Median :0.00  Median :3.000  Median :0.00
Mean   :1.603  Mean   :0.67  Mean   :4.727  Mean   :0.46
3rd Qu.:2.000  3rd Qu.:1.00  3rd Qu.:7.000  3rd Qu.:1.00
Max.   :3.000  Max.   :3.00  Max.   :7.000  Max.   :1.00

```

Figure 5: Represents the summary statistics of the dataframe

From fig.5, we can observe that all the variables has different range of values and we need to standardize our data before proceeding for model building. We can either use the `scale()` function in the `caTools` or use the `preProcess()` function in the `caret` package to standardize the data.

Since the values in our target variable is binary with 0 and 1, it should be a categorical variable. We use the `factor()` function to convert the target variable V14 into factor data structure.

```

> class(training[["v14"]])
[1] "integer"
> training[["v14"]] = factor(training[["v14"]])
> testing[["v14"]] = factor(testing[["v14"]])
> class(training[["v14"]])
[1] "factor"

```

Figure 6: Represents the class of the target variable V14

From Fig.6, we can understand that the class of the target variable V14 has changed from integer to factor.

We can use the `table()` function to check the number of 0's and 1's in the target variable V14

```

> table(heart_df$v14)
 0    1
162 138

```

```

> table(training[["v14"]])
 0    1
112  98

```

From the above, we can interpret that out of the 300 observations, the absence and the presence of the heart disease are 162 and 138, which is 54% and 46% respectively. Similarly, with the training data of 210 observations, the absence and the presence of the heart disease are 112 and 98, which is 53% and 47% respectively. We can say the distribution of the target variable in the training data almost matches with the total dataset.

## **4. Model Building**

The model building is performed using the `train()` function in the `caret` package. This function supports various classification and regression algorithm by setting up sets up a grid of tuning parameters. The fitted model is used for predicting the target variable by inputting predictor values.

To control the computational nuances of the `train()` method, we use the `trainControl()` method that generates parameters that further control how models are created. We are setting 3 parameters of `trainControl()` method. The “method” parameter holds the details about resampling method. We can set “method” with many values like "boot", "cv", "LOOCV", "LGOVCV", "repeatedcv", "timeslice", "none" and "oob". For this project, we use `repeatedcv`, which repeatedly performs X-fold cross-validation on the training data.

The “number” parameter controls with the number of folds in K-fold cross-validation. The “repeats” parameter contains the complete sets of folds to compute for our repeated cross-validation. Since we are using `number = 10` and `repeats = 3`, it will perform 10-fold cross-validation on the training data 3 times, using a different set of folds for each cross-validation. This `trainControl()` methods returns a list which will be pass on our `train()` method.

### **4.1 SVM Classifier using Linear Kernel**

First, let's create a model for SVM linear classifier with the assumption that the training data are linearly separable. A hyperplane should be at a maximum distance from the nearest data point of either class. For training SVM linear classifier, we should pass “`svmLinear`” with “method” parameter in `train()` function. The “`V14~.`” denotes a formula for using all attributes in our classifier and V14 as the target variable. The “`trControl`” parameter holds the list of values that controls the `train()` function. It should be passed with results from our `trainControl()` method.

The data pre-processing can be carried out by passing desired values in the `preProcess` argument in the `train()` function of the `caret` package. The scale transform calculates the standard deviation for every variable and divides each value by its standard deviation. The center transform calculates the mean for every variable and subtracts it from each value. By combining both scale and center transforms, the variables are standardized to a mean value of 0 and a standard deviation of 1. (Nselvar, 2022)

The “`tuneLength`” parameter holds an integer value that represents the number of levels for each tuning parameter that should be generated by `train`. By setting the `tuneLength` parameter as 10, it allows to pick 10 arbitrary values for the `C`, the “cost” of the linear kernel. The `tuneLength` parameter controls the complexity of the boundary between support vectors.

```
> svm_Linear
Support Vector Machines with Linear kernel

210 samples
 13 predictor
  2 classes: '0', '1'

Pre-processing: centered (13), scaled (13)
Resampling: Cross-validated (10 fold, repeated 3 times)
Summary of sample sizes: 189, 188, 189, 188, 190, 189, ...
Resampling results:

   Accuracy   Kappa
0.7759163  0.5482176

Tuning parameter 'C' was held constant at a value of 1
```

*Figure 7: The output of the SVM linear fit model*

From fig.7, we understand that the support vector machine model was created with the linear kernel. The testing data used for creating this model contains 210 samples with 13 predictor variables. The target variable has 2 classes which are ‘0’ and ‘1’. The centre and scale transformation were performed on 13 predictor variables. The resampling was done with 10-fold cross-validation on the training data 3 times, using a different set of folds for each cross-validation. The resulting SVM linear model has an accuracy of 77.59% and the kappa value of 0.5482176. By default, the tuning parameter ‘C’ was held constant at a value of 1. The cost, C is a hyperparameter which is set before the training model and used to control error.

## 4.2 Testing the SVM linear classifier

To test the prediction of the SVM linear model, we use the `predict()` function which is part of the `caret` package. We will be passing 2 arguments. It's first parameter is our trained model 'svm\_Linear' and second parameter "newdata" holds our testing data frame. The `predict()` method returns a list which is stored in a `test_pred` variable. The `test_pred` has the output of "0" and "1" which is the predicted class for the test dataset.

To check the model's test accuracy, we use the `confusionMatrix()` function in R which shows a cross-tabulation of the observed and predicted classes..

```
> confusionMatrix(test_pred, testing$v14)
Confusion Matrix and Statistics

          Reference
Prediction 0  1
          0 40  5
          1 10 35

              Accuracy : 0.8333
              95% CI   : (0.74, 0.9036)
    No Information Rate : 0.5556
    P-Value [Acc > NIR] : 2.25e-08

              Kappa : 0.6667

  Mcnemar's Test P-Value : 0.3017

              Sensitivity : 0.8000
              Specificity : 0.8750
    Pos Pred Value : 0.8889
    Neg Pred Value : 0.7778
        Prevalence : 0.5556
    Detection Rate : 0.4444
    Detection Prevalence : 0.5000
    Balanced Accuracy : 0.8375

    'Positive' Class : 0
```

Figure 8: The confusion matrix output of the test data for SVM linear fit model

From fig.8, we can observe that out of the 90-test observation, the absence and presence of the heart disease is 45 observations each. In the absence of the heart disease class, 89% data are correctly classified whereas in presence of the heart disease class only 78% data are correctly classified. The model accuracy is 83.33%. The 95% confidence interval lies between 0.74 and 0.9036. The "no-information rate" which is the largest proportion of the



observed classes is 0.5556. McNemar's P-value test indicates that the accuracy of the model is higher than NIR. The Kappa value is 0.6667.

### 4.3 Tuning the SVM linear classifier

We can tune the model for better accuracy by using the tuneGrid parameter to do some sensitivity analysis around the values C value(Cost) between 0 and 5 in Linear classifier. This can be done by inputting values in grid search. The next code snippet will show you, building & tuning of an SVM classifier with different values of C. The expand.grid() function is used to build a dataframe that contain all the combinations of C we want to look at. Next step is to use this data frame for testing our classifier at specific C values. It needs to be put in train() method with tuneGrid parameter(Revolutions, 2022).

```
> svm_Linear_Grid
Support Vector Machines with Linear kernel

210 samples
 13 predictor
  2 classes: '0', '1'

Pre-processing: centered (13), scaled (13)
Resampling: Cross-validated (10 fold, repeated 3 times)
Summary of sample sizes: 189, 189, 190, 189, 189, 188, ...
Resampling results across tuning parameters:
```

| C     | Accuracy  | Kappa     |
|-------|-----------|-----------|
| 0.000 | NaN       | NaN       |
| 0.001 | 0.5854690 | 0.1180914 |
| 0.010 | 0.8126623 | 0.6197360 |
| 0.050 | 0.8078211 | 0.6103251 |
| 0.100 | 0.7902020 | 0.5759191 |
| 0.250 | 0.7759091 | 0.5468493 |
| 0.500 | 0.7683405 | 0.5323466 |
| 0.750 | 0.7653824 | 0.5269904 |
| 1.000 | 0.7668182 | 0.5295216 |
| 1.250 | 0.7651515 | 0.5261955 |
| 1.500 | 0.7651587 | 0.5265094 |
| 1.750 | 0.7667460 | 0.5297019 |
| 2.000 | 0.7682612 | 0.5325831 |
| 5.000 | 0.7681890 | 0.5327688 |

```
Accuracy was used to select the optimal model using the largest value.
The final value used for the model was c = 0.01.
```

Figure 9: The output of the SVM linear grid fit model

Since, we have used the dataframe grid to impute the values of C in the range of 0 to 5, we obtain the accuracy and kappa value for each of the cost value. The optimal model is obtained for the cost value with highest accuracy. From fig.9, we can observe that for the C value of 0.01, the maximum accuracy of 81.26% is attained.

The plot of the svm\_Linear\_Grid shows that the maximum accuracy is reached when  $c=0.01$ .




Figure 10: Plot of Cost versus accuracy (repeated cross-validation) for SVM Linear

We tested the prediction model of svm\_Linear\_Grid with our test data using the predict() function.

```
> confusionMatrix(test_pred_grid, testing$V14 )
Confusion Matrix and Statistics

      Reference
Prediction 0  1
      0  47  4
      1   3 36

      Accuracy : 0.9222
      95% CI   : (0.8463, 0.9682)
No Information Rate : 0.5556
P-Value [Acc > NIR] : 1.851e-14

      Kappa : 0.8421

McNemar's Test P-Value : 1

      Sensitivity : 0.9400
      Specificity : 0.9000
      Pos Pred Value : 0.9216
      Neg Pred Value : 0.9231
      Prevalence : 0.5556
      Detection Rate : 0.5222
      Detection Prevalence : 0.5667
      Balanced Accuracy : 0.9200

      'Positive' Class : 0
```

Figure 11: The confusion matrix output of the test data for SVM linear grid fit model

The confusion matrix calculated for SVM Linear Grid model shows an accuracy of 92.22%. we can observe that out of the 90-test observation, the absence and presence of the heart disease are 51 and 39 observations respectively. In both absence and presence of the heart disease class, 92% data are correctly classified. The 95% confidence interval lies between 0.8463 and 0.9682. The Kappa value obtained is 0.8421.

We can observe that the accuracy of the SVM linear model has increased from 83.33% to 92.22% by using the Cost value of 0.01 instead of the default value 1.

## 5. SVM Classifier using Non-Linear Kernel

We will now try to build a SVM classifier model using Non-Linear Kernel like Radial Basis and Polynomial Function. For radial basis function, we draw completely non-linear hyperplanes between the 2 classes. For distance metric, squared euclidean distance is used. For using RBF kernel, we need to pass “svmRadial” as “method” parameter in train() function. Selecting the right Cost “C” and “sigma” parameter is critical for Radial kernel.

```
> svm_Radial
Support Vector Machines with Radial Basis Function Kernel

210 samples
 13 predictor
  2 classes: '0', '1'

Pre-processing: centered (13), scaled (13)
Resampling: Cross-validated (10 fold, repeated 3 times)
Summary of sample sizes: 189, 188, 189, 188, 190, 189, ...
Resampling results across tuning parameters:


   C      Accuracy   Kappa
0.25  0.8170779  0.6288443
0.50  0.8136652  0.6219408
1.00  0.7993723  0.5934927
2.00  0.8025469  0.6007756
4.00  0.7976984  0.5912452
8.00  0.7927850  0.5828339
16.00 0.7946681  0.5868178
32.00 0.7834055  0.5651129
64.00 0.7631241  0.5248845
128.00 0.7486797  0.4947990

Tuning parameter 'sigma' was held constant at a value of 0.04744242
Accuracy was used to select the optimal model using the largest value.
The final values used for the model were sigma = 0.04744242 and C = 0.25.
```

Figure 12: The output of SVM Radial fit model

From fig.12, we can understand that we have used the same training dataset, pre-processing and traincontrol parameters that we used for SVM linear. The SVM RBF kernel calculates the variation and provides the optimal values of cost(C) and sigma. From the above output, we can say the best value of tuning parameter, sigma is 0.04744242 and it was held constant for the various values of the cost ranging from 0.25 to 128. The optimal model is chosen for the cost value whose accuracy is the largest. Since the accuracy of 81.70% is obtained when the cost value is 0.25, we can say the best cost value,  $C = 0.25$ .

We can plot the output of the SVM radial using the plot() function. The cost versus accuracy plot shows when the cost value,  $C = 0.25$ .



*Figure 13: Plot of Cost versus accuracy (repeated cross-validation) for SVM Radial*

We tested the prediction model of SVM Radial model with our test data using the predict() function. Confusion matrix is used to check the model's test accuracy and it shows a cross-tabulation of the observed and predicted classes.

The confusion matrix calculated for SVM Radial model shows an accuracy of 90%. We can observe that out of the 90-test observation, the absence and presence of the heart disease are 49 and 41 observations respectively. In the absence of the heart disease class, 92% data are correctly classified whereas in presence of the heart disease class only 88% data are correctly classified.. The 95% confidence interval lies between 0.8186 and 0.9532. The Kappa value obtained is 0.798.

```

> confusionMatrix(test_pred_Radial, testing$y14 )
Confusion Matrix and Statistics

      Reference
Prediction 0  1
      0 45  4
      1  5 36

      Accuracy : 0.9
      95% CI : (0.8186, 0.9532)
      No Information Rate : 0.5556
      P-Value [Acc > NIR] : 1.162e-12

      Kappa : 0.798

      Mcnemar's Test P-Value : 1

      Sensitivity : 0.9000
      Specificity : 0.9000
      Pos Pred Value : 0.9184
      Neg Pred Value : 0.8780
      Prevalence : 0.5556
      Detection Rate : 0.5000
      Detection Prevalence : 0.5444
      Balanced Accuracy : 0.9000

      'Positive' Class : 0

```

Figure 14: The confusion matrix output of the test data for SVM Radial fit model

We use `train()`'s `tuneGrid` parameter to test and tune our classifier with different values `C` and `sigma` to produce the model with the best accuracy. The `expand.grid()` function is used to build a dataframe contain all the combinations of `C` and `sigma` that we want to check on the model. The variable `grid_radial` dataframe will hold values of `sigma` & `C`. Value of `grid_radial` will be passed to `train()` method's `tuneGrid` parameter.

We have checked the SVM Radial model with below values of cost and sigma.

`sigma = 0,0.01, 0.02, 0.025, 0.03, 0.04, 0.05, 0.06, 0.07,0.08, 0.09, 0.1, 0.25, 0.5, 0.75,0.9),`

`cost = 0,0.01, 0.05, 0.1, 0.25, 0.5, 0.75, 1, 1.5, 2,5`

The optimal model is chosen for the cost ad sigma value that has high accuracy. Based on the below output the final SVM Radial Grid classifier was modelled by using the values `sigma = 0.025` and `C = 0.1`.

```

> svm_Radial_Grid
Support Vector Machines with Radial Basis Function Kernel

210 samples
 13 predictor
  2 classes: '0', '1'

Pre-processing: centered (13), scaled (13)
Resampling: Cross-validated (10 fold, repeated 3 times)
Summary of sample sizes: 189, 189, 188, 189, 189, 190, ...
Resampling results across tuning parameters:

  sigma  C      Accuracy  Kappa
0.000  0.00      NaN      NaN
0.000  0.01  0.5333766  0.00000000
0.000  0.05  0.5333766  0.00000000
0.000  0.10  0.5333766  0.00000000
0.000  0.25  0.5333766  0.00000000
0.000  0.50  0.5333766  0.00000000
0.000  0.75  0.5333766  0.00000000
0.000  1.00  0.5333766  0.00000000
0.000  1.50  0.5333766  0.00000000
0.000  2.00  0.5333766  0.00000000
0.000  5.00  0.5333766  0.00000000
0.010  0.00      NaN      NaN
0.010  0.01  0.5333766  0.00000000
0.010  0.05  0.5333766  0.00000000
0.010  0.10  0.7590693  0.50182781
0.010  0.25  0.8143434  0.62293518
0.010  0.50  0.8157792  0.62624374
0.010  0.75  0.8094300  0.61313348
0.010  1.00  0.8094300  0.61313348

#####
0.025  0.01  0.5333766  0.00000000
0.025  0.05  0.7211111  0.41943791
0.025  0.10  0.8206999  0.63511957
0.025  0.25  0.8142713  0.62303532
0.025  0.50  0.8110967  0.61637360
0.025  0.75  0.8141919  0.62306824
0.025  1.00  0.8126046  0.61972198
0.025  1.50  0.8078427  0.61046109
0.025  2.00  0.7966450  0.58832646
0.025  5.00  0.7966450  0.58983222

#####
Accuracy was used to select the optimal model using the largest value.
The final values used for the model were sigma = 0.025 and C = 0.1.

```

Figure 15: The confusion matrix output of the test data for SVM Radial Grid fit model




Figure 15: Plot of Cost versus accuracy (repeated cross-validation) for SVM Radial Grid model.

In the fig.15, the colour coded lines represent the individual cost value. We can observe that the maximum accuracy is attained for most of the cost value when the sigma is between 0.01 and 0.1. When sigma is 0.0, the accuracy remains low and constant at 0.5333766 for each of the cost value. Similarly, when the cost is 0.05, the accuracy remains constant at 0.5333766 when the sigma value is above 0.2.

```
> confusionMatrix(test_pred_Radial_Grid, testing$V14 )
Confusion Matrix and Statistics

      Reference
Prediction 0  1
0      46  6
1       4 34

      Accuracy : 0.8889
      95% CI   : (0.8051, 0.9454)
No Information Rate : 0.5556
P-value [Acc > NIR] : 7.675e-12

      Kappa : 0.7739
```

Figure 15: The output of confusion matrix for SVM Radial Grid model

## 5.2. SVM Classifier using Polynomial Kernel

Polynomial Kernel is a general representation of kernels with a degree of more than one. For using polynomial kernel, we need to pass “svmPoly” as “method” parameter in train() function.

```
> svm_polynomial
Support Vector Machines with Polynomial kernel

210 samples
 13 predictor
  2 classes: '0', '1'

Pre-processing: centered (13), scaled (13)
Resampling: Cross-validated (10 fold, repeated 3 times)
Summary of sample sizes: 189, 189, 188, 189, 189, 190, ...
Resampling results across tuning parameters:
```

| degree | scale | C    | Accuracy  | Kappa     |
|--------|-------|------|-----------|-----------|
| 1      | 1e-03 | 0.25 | 0.5333766 | 0.0000000 |
| 1      | 1e-03 | 0.50 | 0.5333766 | 0.0000000 |
| 1      | 1e-03 | 1.00 | 0.5845310 | 0.1168046 |
| 1      | 1e-03 | 2.00 | 0.8033045 | 0.5975437 |
| 1      | 1e-03 | 4.00 | 0.8190332 | 0.6322478 |
| 1      | 1e-02 | 0.25 | 0.8112410 | 0.6143898 |
| 1      | 1e-02 | 0.50 | 0.8142641 | 0.6229704 |
| 1      | 1e-02 | 1.00 | 0.8126768 | 0.6197440 |
| 1      | 1e-02 | 2.00 | 0.8078427 | 0.6100852 |
| 1      | 1e-02 | 4.00 | 0.8078427 | 0.6105737 |
| 1      | 1e-01 | 0.25 | 0.8077633 | 0.6101548 |

```
#####

3      1e-02  1.00  0.8141126  0.6230503
3      1e-02  2.00  0.8014935  0.5979363
3      1e-02  4.00  0.7997403  0.5946458
3      1e-01  0.25  0.7842496  0.5649528
3      1e-01  0.50  0.7860606  0.5684708
3      1e-01  1.00  0.7843867  0.5669926
3      1e-01  2.00  0.7682107  0.5346278
3      1e-01  4.00  0.7399134  0.4777425
3      1e+00  0.25  0.7477128  0.4942010
3      1e+00  0.50  0.7477128  0.4942010
3      1e+00  1.00  0.7477128  0.4942010
3      1e+00  2.00  0.7477128  0.4942010
3      1e+00  4.00  0.7477128  0.4942010
3      1e+01  0.25  0.7476335  0.4947378
3      1e+01  0.50  0.7476335  0.4947378
3      1e+01  1.00  0.7476335  0.4947378
3      1e+01  2.00  0.7476335  0.4947378
3      1e+01  4.00  0.7476335  0.4947378

Accuracy was used to select the optimal model using the largest value.
The final values used for the model were degree = 3, scale = 0.001 and c = 1.
```

The final SVM Polynomial classifier is modelled using the values degree = 3, scale = 0.001 and C = 1.



```

> confusionMatrix(test_pred_polynomial, testing$V14 )
Confusion Matrix and Statistics

      Reference
Prediction 0  1
0      48   6
1       2  34

      Accuracy : 0.9111
      95% CI   : (0.8323, 0.9608)
No Information Rate : 0.5556
P-Value [Acc > NIR] : 1.564e-13

      Kappa : 0.8182

```

Figure 16: The output of confusion matrix for SVM Polynomial model

The confusion matrix of SVM Polynomial model shows the accuracy of 91.11%.

### 5.3. Comparison between SVM models

Let's compare the results of the SVM Linear, Radial and Polynomial models. The test accuracy of each model after tuning is tabled below

| SVM Model      | Accuracy in % |
|----------------|---------------|
| SVM Linear     | 92.22%        |
| SVM Radial     | 90%           |
| SVM Polynomial | 91.11%        |

We can understand that for the given heart disease dataset, SVM Linear classifier can be used as a best prediction model since it has the accuracy of 92.22%.

## 6. Conclusion

Based on the model output, we can say the support vector machine (SVM) is a promising classifier approach for predicting the presence of heart disease with the given prediction variables. Majority of the variables in the dataset contribute significantly for building of a SVM predictive model with good accuracy rate ( $\geq 90\%$ ). SVM Linear classifier has the highest accuracy when compared to Radial and Polynomial classifier which may be due to overfitting.

## References

1. The 5th Tribe, Support Vector Machines and caret. (n.d.). Revolutions. Retrieved June 22, 2022, from <https://blog.revolutionanalytics.com/2015/10/the-5th-tribe-support-vector-machines-and-caret.html>
2. Bank Telemarketing NN & SVM. (n.d.). Nselvar.github.io. Retrieved June 22, 2022, from <https://nselvar.github.io/bank-telemarketing/>